Genomic Data Science · DATA405/BIOL480 · Fall 2026 · Howard University

Lab 1 · Python and R on real chr22 data

Wednesday Oct 7, 2026 · 5:10–6:40 pm · 1.5 hours · Instructor: Jinrui Xu

Name: _____ Lab partner (optional): _____

Before you start

- Open the course website in any browser (Chrome, Edge, Firefox, Safari) on Windows, Mac or Chromebook and go to Lab 1. Nothing is installed; Python and R run inside the page. The first cell in each language takes 10–30 s to start.
- Do every exercise in **both** languages. Predict the output before you press Run. Use the Hint before the Solution.
- ChatGPT is allowed (Exercise 6 requires it). Copilot-style autocomplete is not used in this lab: the exercises are the exact lines it would type for you.
- Write your answers to Q1–Q7 on this handout or in your lab report. Report due Tuesday, Oct 13 on Canvas (one page is enough).

Data used today (all real)

File	What	Used in
`COMT_cds.fa`	coding sequence of COMT (chr22), 816 nt, GRCh38	Ex 2
`reads_5k.fastq`	5,000 real ENCODE RNA-seq reads, 100 nt	Ex 3
`chr22_genes.tsv`	1,747 GENCODE genes on chr22	Ex 4
`chr22_expression.tsv`	RSEM expression of those genes in one sample (BIOL419 Lab 4 pipeline)	Ex 4
`chr22_mismatches.tsv`	read-vs-reference mismatches from that sample	Ex 4
`airway_counts.tsv`, `airway_metadata.tsv`	8-sample RNA-seq count matrix, dexamethasone vs control (Himes 2014)	Ex 5

Exercise 1 · Hello, two languages (10 min)

Run both cells, then change the gene name and the number and run again.

Python

```
gene = "COMT"
n_aa = 271
print(gene, "encodes", n_aa, "amino acids")
print("That needs", n_aa * 3, "nt + a stop")
```

R

```
gene <- "COMT"
n_aa <- 271
cat(gene, "encodes", n_aa, "amino acids\n")
cat("That needs", n_aa * 3, "nt + a stop\n")
```

Q1. In the R cell, what happens if you write n_aa = 271 instead of n_aa <- 271? And in Python, n_aa <- 271? Try both; one sentence.

Exercise 2 · DNA as a string: COMT and rs4680 (20 min)

Run the setup cell first (defines read_fasta and reads the sequence into cds).

Python

```
def read_fasta(path):
    seq = []
    for line in open(path):
        if not line.startswith(">"):
            seq.append(line.strip())
    return "".join(seq)

cds = read_fasta("data/COMT_cds.fa")
print(len(cds), "nt")
```

R

```
read_fasta <- function(path) {
  lines <- readLines(path)
  paste(lines[!startsWith(lines, ">")],
        collapse = "")
}

cds <- read_fasta("data/COMT_cds.fa")
cat(nchar(cds), "nt\n")
```

2a · GC content — fill in the blanks

Python

```
g = cds.count("G")
c = ______
gc = (g + c) / ______
print(round(gc * 100, 1), "%")
```

R

```
bases <- strsplit(cds, "")[[1]]
gc <- mean(bases %in% ______)
cat(round(gc * 100, 1), "%\n")
```

Checkpoint Both languages must print 61.8 %. (The human genome averages about 41 %.)

2b · Reverse complement

Python

```
complement = {"A": "T", "T": "A",
              "G": "C", "C": "G"}
rc = "".join(complement[b] for b in ______)
print(rc[:20])
```

R

```
complement <- c(A = "T", T = "A",
                G = "C", C = "G")
rc <- paste(______(complement[bases]),
            collapse = "")
cat(substr(rc, 1, 20), "\n")
```

Checkpoint The reverse complement must start with TCA, because the CDS ends with the stop codon TGA.

2c · Translate, then mutate codon 158

Run the setup cell on the website (codon table + translate). Then complete:

Python

```
i = (158 - 1) * 3
codon = cds[i:i+3]
print("codon 158:", codon, "->",
codon_table[codon])
mutant = cds[:i] + ______ + cds[i+1:]
print("mutant   :", mutant[i:i+3], "->",
      translate(mutant)[______])
```

R

```
i <- (158 - 1) * 3 + 1
codon <- substr(cds, i, i + 2)
cat("codon 158:", codon, "->",
codon_table[codon], "\n")
mutant <- cds
substr(mutant, i, i) <- ______
cat("mutant   :", substr(mutant, i, i + 2),
"->",
    substr(translate(mutant), ______,
______), "\n")
```

Checkpoint Protein starts MPEAPPLLLA, 271 aa. Codon 158: GTG (V) → ATG (M). This is Val158Met (rs4680).

Q2. Delete the letter instead of substituting it. How many amino acids still match before the sequence goes wrong, and where is the first stop codon? Why is a deletion usually worse than a substitution?

Exercise 3 · Reads: FASTQ and quality (15 min)

Run the setup cell (reads all 5,000 records into seqs and quals).

Python

```
seqs, quals = [], []
with open("data/reads_5k.fastq") as f:
    while True:
        name = f.readline()
        if not name:
            break
        seqs.append(f.readline().strip())
        f.readline()
        quals.append(f.readline().strip())
```

R

```
lines <- readLines("data/reads_5k.fastq")
n <- length(lines)
seqs  <- lines[seq(2, n, by = 4)]
quals <- lines[seq(4, n, by = 4)]
```

3a · Decode quality and plot the mean per position

Python

```
L = len(seqs[0])
mean_q = []
for i in range(L):
    tot = sum(ord(q[i]) - ______ for q in
quals)
    mean_q.append(tot / len(quals))
fig, ax = plt.subplots(figsize=(4, 2.6))
ax.plot(range(1, L + 1), mean_q)
```

R

```
Q <- t(sapply(quals,
       function(s) utf8ToInt(s) - ______))
mean_q <- colMeans(Q)
plot(mean_q, type = "l",
     xlab = "position", ylab = "mean Q")
```

Checkpoint Mean Q ≈ 33 at position 1, ≈ 37 in the middle, ≈ 35 at position 100.

3b · Count reads by content

Python

```
n_N = sum("N" in s for s in seqs)
n_polyA = sum(______ in s for s in seqs)
gc = [(s.count("G") + s.count("C")) / len(s)
      for s in seqs]
print(n_N, n_polyA, round(sum(gc)/len(gc),
3))
```

R

```
n_N <- sum(grepl("N", seqs))
n_polyA <- sum(grepl(______, seqs))
gc <- sapply(strsplit(seqs, ""),
  function(s) mean(s %in% c("G", "C")))
cat(n_N, n_polyA, round(mean(gc), 3), "\n")
```

Checkpoint 6 reads contain N; 63 contain 15 A's in a row; mean GC ≈ 0.474.

Q3. The reads are ~47 % GC, the genome is 41 %, COMT is 62 %. Give one biological reason RNA-seq reads are GC-richer than the genome.

Exercise 4 · Tables: genes, expression, mismatches (20 min)

Python

```
import pandas as pd
genes = pd.read_csv("data/chr22_genes.tsv",
sep="\t")
expr  =
pd.read_csv("data/chr22_expression.tsv",
sep="\t")
mm    =
pd.read_csv("data/chr22_mismatches.tsv",
sep="\t")
```

R

```
genes <- read.delim("data/chr22_genes.tsv")
expr  <-
read.delim("data/chr22_expression.tsv")
mm    <-
read.delim("data/chr22_mismatches.tsv")
```

4a · Filter, sort, count

How many protein-coding genes are on the minus strand? Which three protein-coding genes are longest?

Python

```
pc = genes[genes["gene_type"] ==
"protein_coding"]
print((pc["strand"] == ______).sum())
print(pc.sort_values("length_bp",
      ascending=______)[["gene_name",
"length_bp"]].head(3))
```

R

```
pc <- genes[genes$gene_type ==
"protein_coding", ]
sum(pc$strand == ______)
head(pc[order(______pc$length_bp),
        c("gene_name", "length_bp")], 3)
```

4b · Join and summarize

Join genes to expr by gene_id; for each gene_type compute the fraction of genes with TPM > 1.

Python

```
m = genes.merge(expr[["gene_id", "TPM"]],
on=______)
frac = m.groupby("gene_type")["TPM"] \
        .apply(lambda s: (s > 1).mean())
print(frac.sort_values(ascending=False).head
(5))
```

R

```
m <- merge(genes, expr[, c("gene_id",
"TPM")],
           by = ______)
frac <- tapply(m$TPM > 1, m$gene_type, mean)
head(sort(frac, decreasing = TRUE), 5)
```

> **Checkpoint** About 47 % of protein-coding genes have TPM > 1. The top five expressed genes are immunoglobulin genes and pseudogenes.

Q4. What cell type do the top-expressed genes suggest? Why can a pseudogene (ACTBP15) show a high TPM even though it is not transcribed? (Hint: what is ACTBP15 a copy of, and what was in the reference?)

4c · Mismatches: which look like real variants?

Filter to depth ≥ 10 and allele fraction between 0.3 and 0.7 (what a heterozygous SNP looks like). Then tabulate the most common substitution types.

Python

```
het = mm[(mm["depth"] >= 10) &
         (mm["alt_fraction"] >= ______) &
         (mm["alt_fraction"] <= ______)]
print(len(het))
subst = mm["ref"] + ">" + mm["alt"]
print(subst.value_counts().head(4))
```

R

```
het <- mm[mm$depth >= 10 &
          mm$alt_fraction >= ______ &
          mm$alt_fraction <= ______, ]
nrow(het)
subst <- paste0(mm$ref, ">", mm$alt)
head(sort(table(subst), decreasing = TRUE),
4)
```

Checkpoint Only 8 of 1,097 sites look heterozygous. The two most common substitutions are T>C and A>G.

Q5. A>G (and its mirror T>C) dominating RNA-seq mismatches is a known biological signal, not an error. What is it, and why does DNA sequencing not show it?

Exercise 5 · Statistics on counts: the t-test and its limits (15 min)

Python

```
import numpy as np
from scipy import stats
cts  = pd.read_csv("data/airway_counts.tsv",
                   sep="\t", index_col=0)
meta =
pd.read_csv("data/airway_metadata.tsv",
sep="\t")
logc = np.log2(cts + 1)
ctrl = list(meta["id"][meta["dex"] ==
"control"])
trt  = list(meta["id"][meta["dex"] ==
"treated"])
```

R

```
cts  <- read.delim("data/airway_counts.tsv",
                   row.names = 1)
meta <-
read.delim("data/airway_metadata.tsv")
logc <- log2(cts + 1)
ctrl <- meta$id[meta$dex == "control"]
trt  <- meta$id[meta$dex == "treated"]
```

5a · One gene (ENSG00000152583, SPARCL1): boxplot and t-test

Python

```
g = "ENSG00000152583"
a = logc.loc[g, ctrl]
b = logc.loc[g, trt]
fig, ax = plt.subplots(figsize=(3, 3))
ax.boxplot([a, b], labels=["control",
"dex"])
print(stats.ttest_ind(______, ______))
```

R

```
g <- "ENSG00000152583"
a <- unlist(logc[g, ctrl])
b <- unlist(logc[g, trt])
boxplot(list(control = a, dex = b))
t.test(______, ______)
```

5b · The 19,271-tests trap

Python

```
p = stats.ttest_ind(logc[trt], logc[ctrl],
axis=1).pvalue
print("p < 0.05:", (p < ______).sum(), "of",
len(p))
rng = np.random.default_rng(1)
cols = list(cts.columns); rng.shuffle(cols)
p_shuf = stats.ttest_ind(logc[cols[:4]],
                         logc[cols[4:]],
axis=1).pvalue
print("shuffled:", (p_shuf < 0.05).sum())
```

R

```
tt <- function(a, b) tryCatch(t.test(a, b)
$p.value,
                              error =
function(e) NA)
p <- apply(logc, 1, function(x) tt(x[trt],
x[ctrl]))
cat("p < 0.05:", sum(p < ______, na.rm =
TRUE), "\n")
set.seed(2)
cols <- sample(colnames(cts))
p_shuf <- apply(logc, 1,
  function(x) tt(x[cols[1:4]],
x[cols[5:8]]))
cat("shuffled:", sum(p_shuf < 0.05, na.rm =
TRUE), "\n")
```

> **Checkpoint** With random labels, hundreds to over a thousand genes still pass p < 0.05. Now apply Benjamini–Hochberg (p.adjust(p, "BH") in R) — how many survive at FDR 0.05?

Q6. Two reasons the t-test is the wrong tool here: what does it assume about the distribution of counts, and what does it ignore about the four donors?

Exercise 6 · The apprentice: catch ChatGPT's bug (10 min)

- Ask ChatGPT (TRACE prompt): "Write a Python (then R) function reverse_complement(dna) for DNA sequences that may contain N."
- Paste its function into the Exercise 6 cell on the website and run the three tests: ATGC → GCAT, AAANTT → AANTTT, atgc → GCAT.
- If any test fails, fix the function yourself. Paste the fixed version (both languages) in your report with the ChatGPT thread link.

Q7. The typical first answer fails two of the three tests. Name the two bugs. Run it in both languages: in which language does it crash loudly, and in which does it silently give a wrong answer? Which is worse, and why?

Lab report (due Tuesday, Oct 13)

Answers to Q1–Q7 · the fixed reverse_complement in both languages · one figure from Exercise 3 or 5 exported at 3.2 × 3.2 inches (from the local setup, or right-click → save image from the website). One page is enough.

Data provenance and local setup (Windows/Mac): see the Data and Setup pages of the course website.